



Intro to Databases

SQL

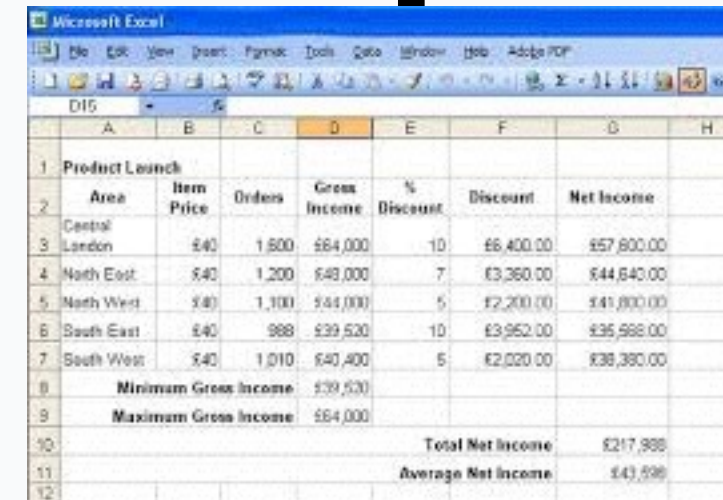
What is a database?

Things that hold info

Accessible




Persistent

A screenshot of a Microsoft Excel spreadsheet showing a product launch analysis. The spreadsheet includes columns for Product Launch, Area, Item Price, Orders, Gross Income, % Discount, and Net Income. It lists data for various regions like Central, London, North East, North West, South East, and South West, along with summary rows for Minimum Gross Income, Maximum Gross Income, Total Net Income, and Average Net Income.

	A	B	C	D	E	F	G	H
1	Product Launch							
2	Area	Item Price	Orders	Gross Income	% Discount	Discount	Net Income	
3	Central	\$40	1,800	\$64,000	10	\$6,400.00	\$57,600.00	
4	London	\$40	1,200	\$48,000	7	\$3,360.00	\$44,640.00	
5	North East	\$40	1,100	\$44,000	5	\$2,200.00	\$41,800.00	
6	North West	\$40	988	\$39,520	10	\$3,952.00	\$35,568.00	
7	South East	\$40	1,010	\$40,400	5	\$2,020.00	\$38,380.00	
8	South West							
9	Minimum Gross Income			\$39,520				
10	Maximum Gross Income			\$64,000				
11						Total Net Income	\$217,988	
12						Average Net Income	\$43,598	

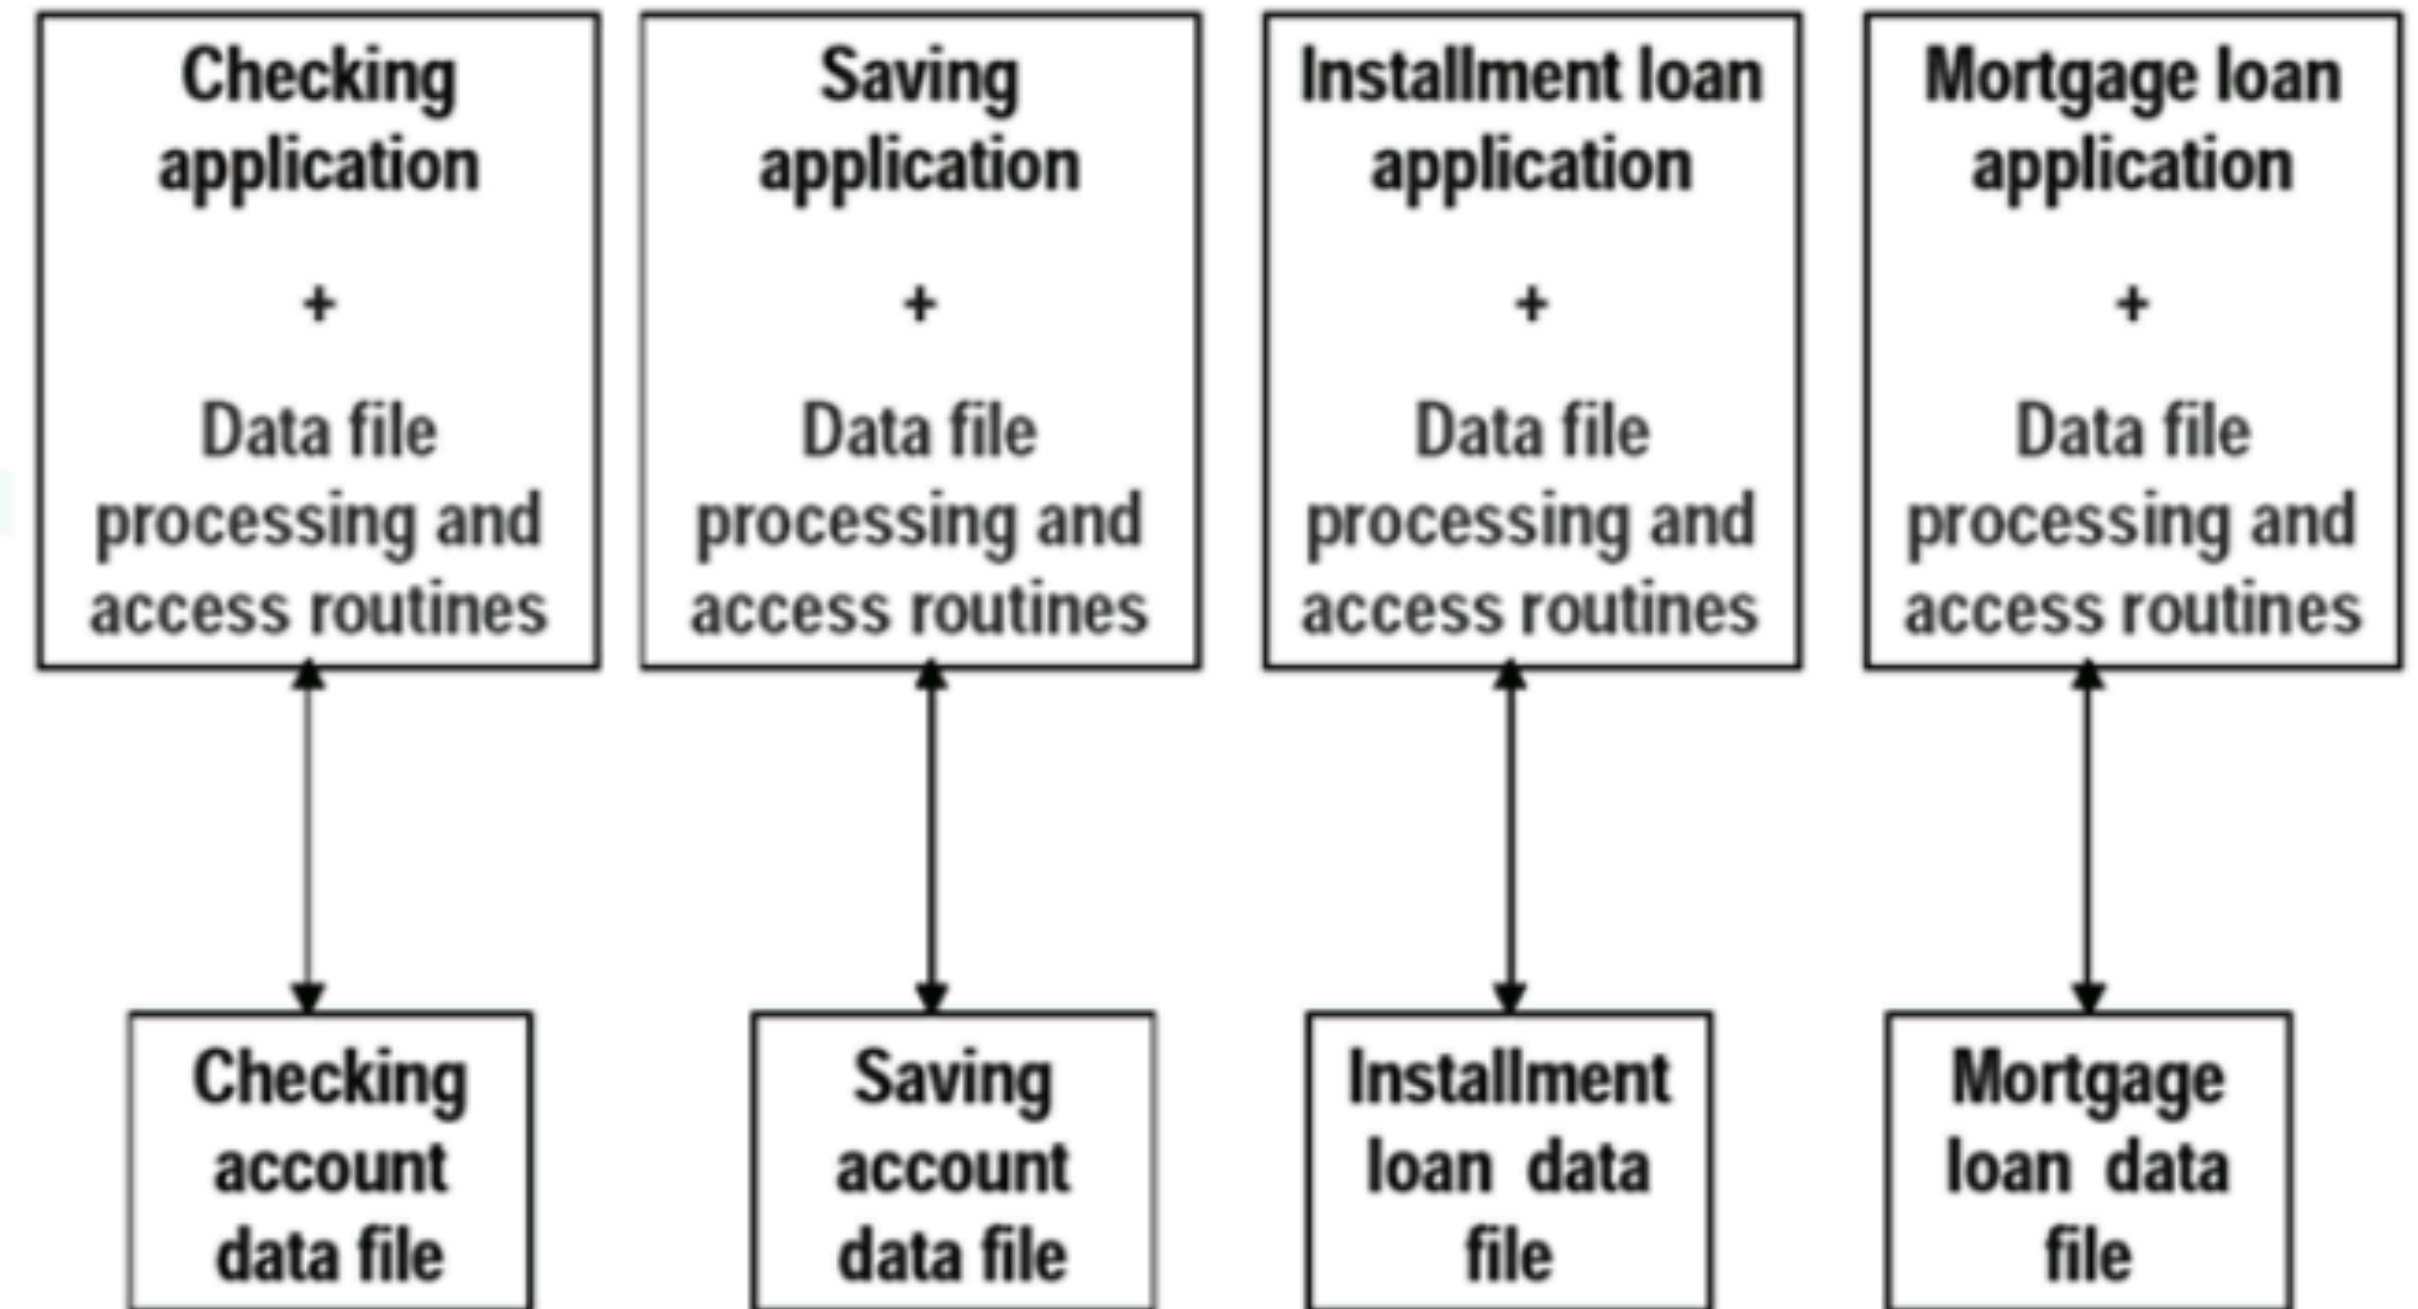


A database **persists** information and is **accessible**
via code

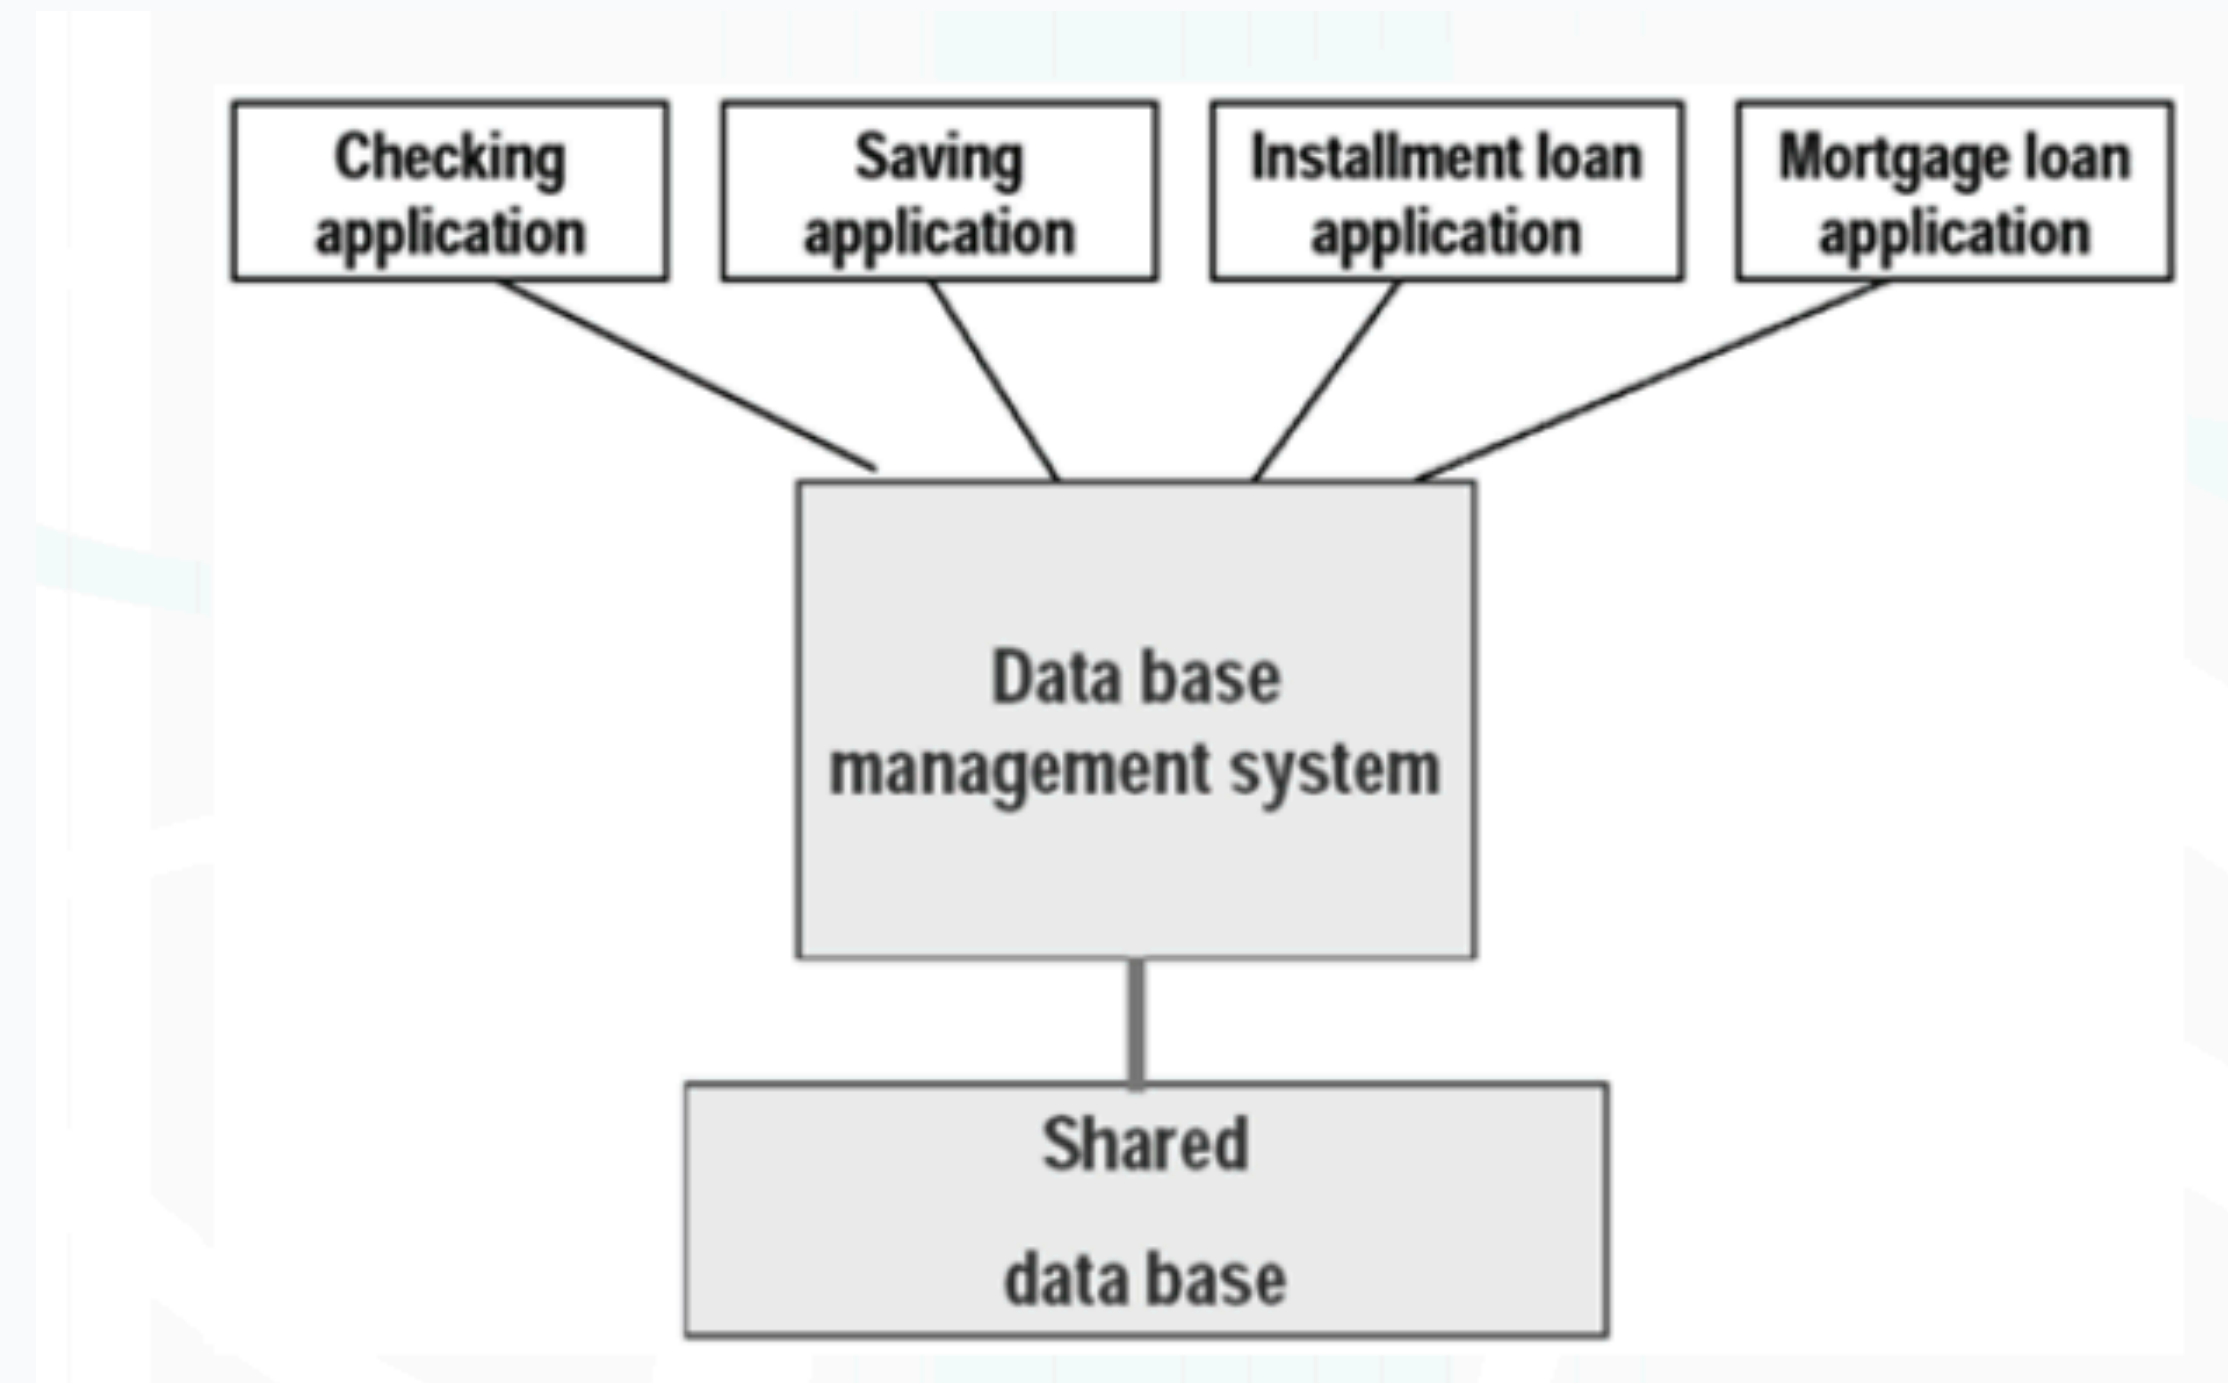


Before Relational DBs (ca. < 1970s)

- Data stored in custom “data files”
- Queried via application-specific code
- Advantages
 - Middle layer not needed
 - Solutions customized for each application
- Disadvantages
 - Hard to change the system
 - Knowledge not compounding
 - Data-transfer is difficult



Database Management Systems (DBMS)




- One layer and language to store and access data
- Sold as a way for "non-technical people" to manage data

“Future users of large data banks must be protected from having to know how the data is organized in the machine (the internal representation).”

-E.F. Codd,

A Relational Model of Data for Large Shared Data Banks

Relational Databases & Logic



- 1969: Edgar Frank "Ted" Codd outlines *relational model* of data
- Wrote Alpha (never implemented) as a *query language*
- IBM slow to adopt his ideas
 - Competitors started to do so
 - IBM team formed without Codd, created **Structured English Query Lang**
- SEQUEL way better than what came before
 - 1979: copied by Larry Ellison (from pre-launch papers / talks!) as "SQL"
- SQL became the standard (ANSI 1986, ISO 1987)
 - Codd continued to fault SQL compared to his theoretical model
 - The Third Manifesto: solve the *object-relational impedance mismatch*



Oracle

Ed Oates, Bruce Scott, Bob Miner, Larry Ellison

Appreciating Databases

- Ubiquitous
- Standardized
- Complex / deep
- Powerful: database admins are...
 - ...feared by developers.
 - ...but also taken for granted until things break.
 - ...befriended by business people.
 - ...contacted by the government for secret data (e.g. NSA).

RDBMS

- Data is stored in relations (tables)
- A simple, structured query language: SQL
 - Programmers can specify what answers a query should return, but not how the query is executed or where and how the data is stored
 - DBMS picks an execution strategy based on indexes, data, workload etc.
- Multi-user, Multi-threaded
 - Multiple processes can access database at same time

Definitions in a Database

- DBs are a collection of Tables (or relations)
- Tables have Columns (attributes) and Rows (instances or tuples)
- Duplicate rows are not allowed
- Rows often have a primary key (ID)

Schema and Content

- Schema: table's blueprint for data shape/format
- Content: actual data (a row) e.g. {1, "Bart S.", 10, "M"}
- A schema is used to validate incoming content

SQL

SQL is used to create/read/update/delete (CRUD) data from a database

- INSERT: Insert new rows into a table
- SELECT: The SELECT command is used to get data from a database
- UPDATE: Update existing rows in a table
- DELETE: Delete rows from a table
- (bonus) CREATE: Make new tables/views/indexes

Example DB

Student

ID	Name	Age	Gender
1	Bart S.	10	M
2	Lisa S.	8	F
3	Jim F.	13	M
4	Joan B.	15	F

Enrollment

StudentID	SchoolID
1	1
2	1
3	2
4	3

School

ID	Name	Level
1	Springfield Elementary	E
2	Brook Middle	M
3	Springbrook High	H
4	Simpson Univ	U

SQL by Example – Select

Student

ID	Name	Age	Gender
1	Bart S.	10	M
2	Lisa S.	8	F
3	Jim F.	13	M
4	Joan B.	15	F

SELECT *
FROM Student
WHERE age > 12

ID	Name	Age	Gender
3	Jim F.	13	M
4	Joan B.	15	F



An SQL query walks into a bar and sees two tables.
He walks up to them and says "Can I join you?"

<https://lol.browserling.com/tables.png>

A more interesting select

Student

ID	Name	Age	Gender
1	Bart S.	10	M
2	Lisa S.	8	F
3	Jim F.	13	M
4	Joan B.	15	F

Enrollment

StudentID	SchoolID
1	1
2	1
3	2
4	3

School

ID	Name	Level
1	Springfield Elementary	E
2	Brook Middle	M
3	Springbrook High	H
4	Simpson Univ	U

Let's say we want to find **all students from Springfield Elementary**. The student table doesn't list the school. We have to use the enrollment table. Will this take two steps?

A more interesting select

Student

ID	Name	Age	Gender
1	Bart S.	10	M
2	Lisa S.	8	F
3	Jim F.	13	M
4	Joan B.	15	F

Enrollment

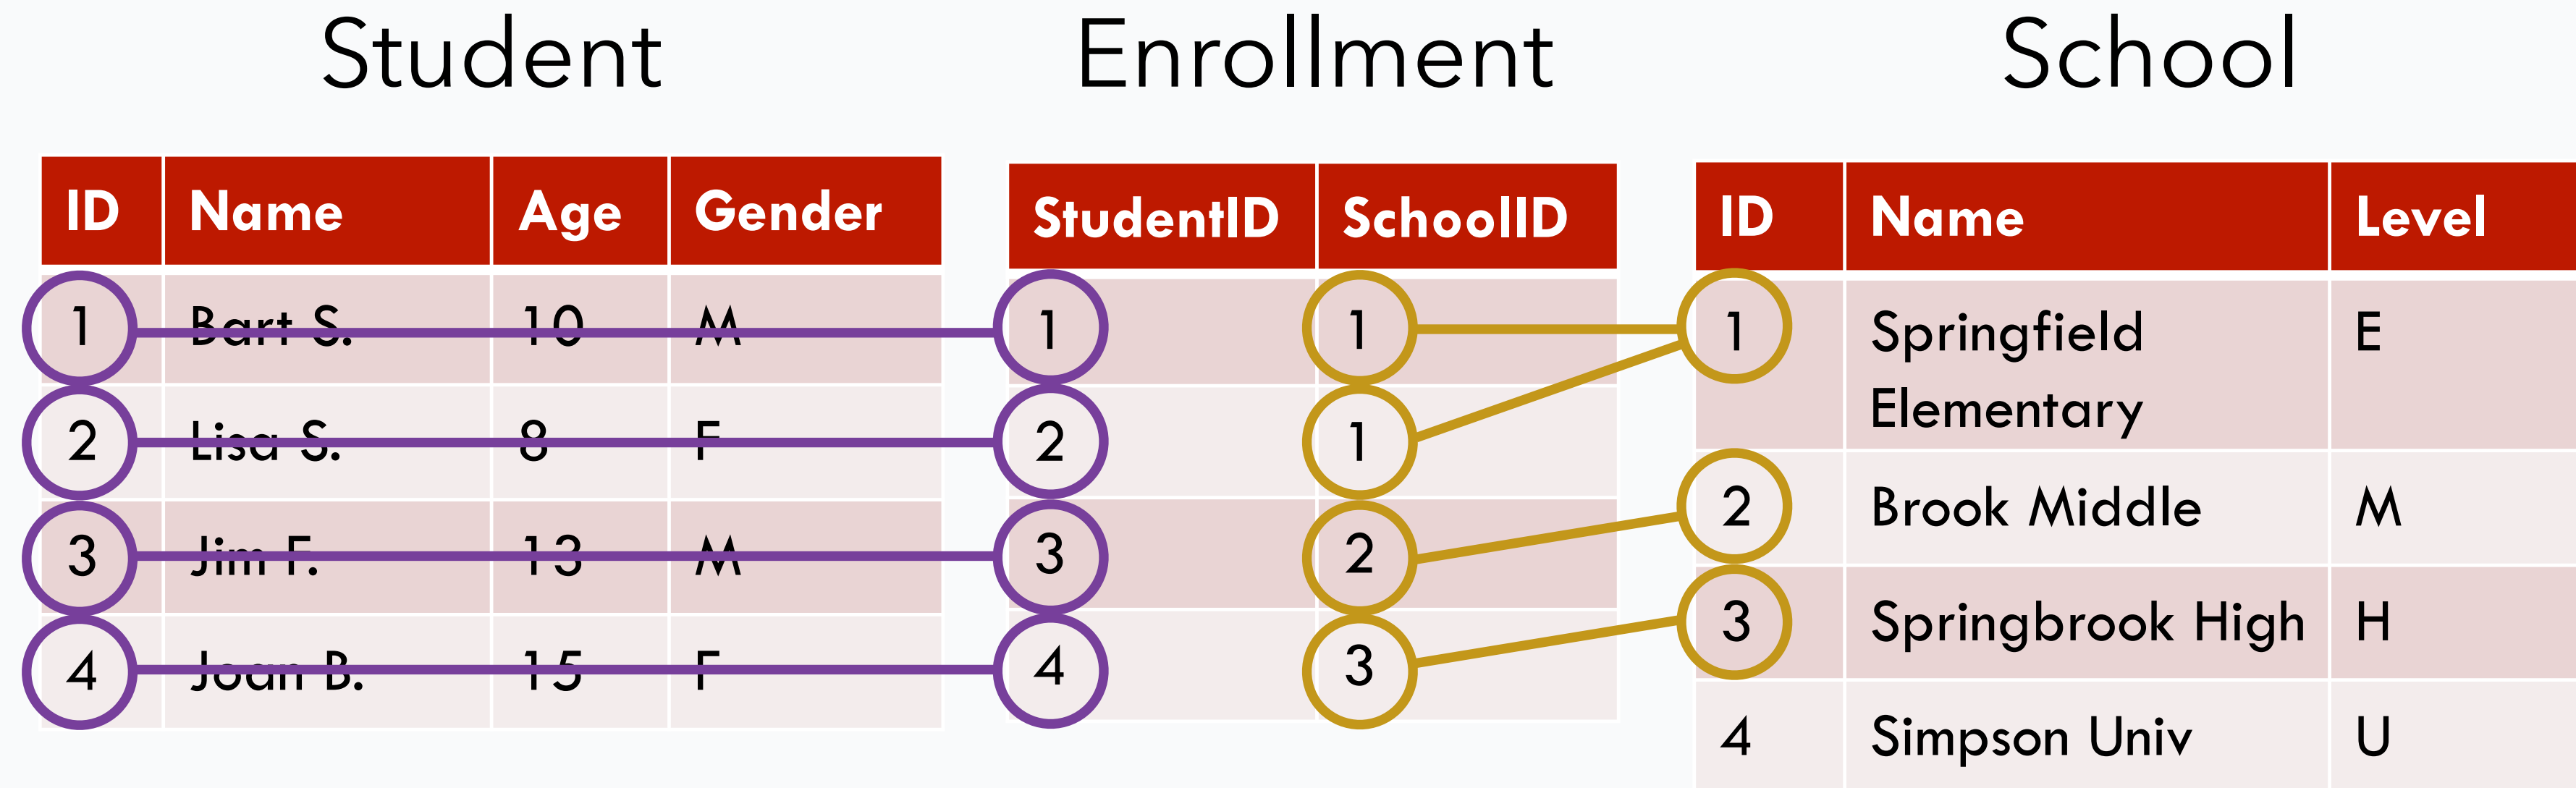
StudentID	SchoolID
1	1
2	1
3	2
4	3

School

ID	Name	Level
1	Springfield Elementary	E
2	Brook Middle	M
3	Springbrook High	H
4	Simpson Univ	U

In fact, we can find all the students from Springfield Elementary (ID: 1) in one SQL statement using a JOIN
A SQL JOIN is used to combine rows from two or more tables, based on a common field between them.
Can you visualize it?

A more interesting select



In fact, we can find all the students from Springfield Elementary (ID: 1) in one SQL statement using a JOIN. A SQL JOIN is used to combine rows from two or more tables, based on a common field between them. Can you visualize it?

SQL Joining

If we joined the **Student** and **School** tables using the data in the Enrollment table, here is how it could look:

Student ID	Name	Age	Gender	School ID	School Name	Level
1	Bart S.	10	M	1	Springfield Elementary	E
2	Lisa S.	8	F	1	Springfield Elementary	E
3	Jim F.	13	M	2	Brook Middle	M
4	Joan B.	15	F	3	Springbrook High	H

```
SELECT *  
FROM
```

```
Student INNER JOIN Enrollment ON Student.id = Enrollment.StudentID  
INNER JOIN School ON Enrollment.SchoolID = School.id
```

SQL Joining


If we joined the **Student** and **School** tables using the data in the Enrollment table, here is how it could look:

Student ID	Name	Age	Gender	School ID	School Name	Level
1	Bart S.	10	M	1	Springfield Elementary	E
2	Lisa S.	8	F	1	Springfield Elementary	E
3	Jim F.	13	M	2	Brook Middle	M
4	Joan B.	15	F	3	Springbrook High	H

```
SELECT *  
FROM  
    Student INNER JOIN Enrollment ON Student.id = Enrollment.StudentID  
    INNER JOIN School ON Enrollment.SchoolID = School.id  
WHERE Enrollment.SchoolID = 1
```


Inner Join

```
SELECT *  
FROM A  
INNER JOIN B  
ON A.Key = B.Key
```




Outer Join

```
SELECT *  
FROM A  
FULL OUTER JOIN B  
ON A.Key = B.Key
```




Left Join

```
SELECT *  
FROM A  
LEFT JOIN B  
ON A.Key = B.Key
```




Right Join

```
SELECT *  
FROM A  
RIGHT JOIN B  
ON A.Key = B.Key
```



<http://www.codeproject.com/Articles/33052/Visual-Representation-of-SQL-Joins>

Inner Join



```
SELECT pets.name, owners.name  
FROM owners  
INNER JOIN pets  
ON pets.OwnerID = owners.ID
```

OWNERS

ID	name
1	Geordi
2	Janeway
3	Data
4	Spock

PETS


ID	ownerID	type	name
1	4	Monkey	Mittens
2	null	Lizard	Carol
3	1	Dog	Rufus
4	2	Cat	Fireball

pets.name	owners.name
Mittens	Spock
Rufus	Geordi
Fireball	Janeway


PETS

ID	ownerID	type	name
1	4	Monkey	Mittens
2	null	Lizard	Carol
3	1	Dog	Rufus
4	2	Cat	Fireball

pets.name	owners.name
Mittens	Spock
Rufus	Geordi
Fireball	Janeway
null	Data



Left Join




```
SELECT pets.name, owners.name  
FROM owners  
LEFT JOIN pets  
ON pets.OwnerID = owners.ID
```

OWNERS

ID	name
1	Geordi
2	Janeway
3	Data
4	Spock

PETS




pets.name	owners.name
Mittens	Spock
Carol	null
Rufus	Geordi
Fireball	Janeway

ID	ownerID	type	name
1	4	Monkey	Mittens
2	null	Lizard	Carol
3	1	Dog	Rufus
4	2	Cat	Fireball

OWNERS

ID	name
1	Geordi
2	Janeway
3	Data
4	Spock

Right Join




```
SELECT pets.name, owners.name  
FROM owners  
RIGHT JOIN pets  
ON pets.OwnerID = owners.ID
```


OWNERS

ID	name
1	Geordi
2	Janeway
3	Data
4	Spock

Outer Join



```
SELECT pets.name, owners.name  
FROM owners  
FULL OUTER JOIN pets  
ON pets.OwnerID = owners.ID
```

PETS

	pets.name	owners.name
	Mittens	Spock
➔	Carol	null
	Rufus	Geordi
	Fireball	Janeway
➔	null	Data

ID	ownerID	type	name
1	4	Monkey	Mittens
2	null	Lizard	Carol
3	1	Dog	Rufus
4	2	Cat	Fireball

WORKSHOP